

# From UDEs to PINNs

---



# UDEs and Similar Approaches

**UDEs provide a good starting point for two more approaches**

If you keep the connection to physics, but you relax the ODE mechanism

- ...Then you get **Physics Informed Neural Networks**
  - Technically, UDEs can be considered PINNs
  - ...But the term refers typically to the approaches surveyed (e.g.) [here](#)

If you keep the ODE mechanism, but you drop the connection to physics

- ...Then you get **Neural Ordinary Differential Equations**
  - This was the first approach to integrate NNs and differential equations
  - The seminal paper is [publicly available](#)

**We are going to briefly outline the former approach**



# From UDEs to PINNs

## Let's start by recapping how UDEs work

At **inference time**, we (typically) integrate an initial value problem:

$$\begin{aligned}\dot{\hat{y}} &= f(\hat{y}, t, U(\hat{y}, t, \theta)) \\ \hat{y}(0) &= y_0\end{aligned}$$

At **training time**, we solve:

$$\begin{aligned}\operatorname{argmin}_{\theta} \quad & L(\hat{y}(t), y) \\ \dot{\hat{y}} &= f(\hat{y}, t, U(\hat{y}, t, \theta)) \\ \hat{y}(t_0) &= y_0\end{aligned}$$

- Which requires to embed ODE integration in gradient descent



# From UDEs to PINNs

**What if we tried to simplify the inference process?**

For example, we could use a NN to approximate  $y(t)$  itself

$$\hat{y}(t; \theta) \simeq y(t)$$

**This approach has several immediate benefits:**

- Inference becomes as efficient as evaluating  $\hat{y}(t; \theta)$ 
  - No need to integrate anything, linear scalability w.r.t. the sampling points
- Handling PDEs also becomes pretty simple
  - We just need to use a multivariate  $t$

**But where is physics here?**



# Training in PINNs

The ODE is taken into account **at training time**

Superficially, the training problem is similar to the UDE one:

$$\begin{aligned}\operatorname{argmin}_{\theta} L(\hat{y}(t, \theta), y) \\ \dot{\hat{y}}(t; \theta) &= f(\hat{y}(t; \theta), t) \\ \hat{y}(t_0; \theta) &= y_0\end{aligned}$$

...But in fact, the situation is very different:

- Since both  $\hat{y}(t; \theta)$  and  $\dot{\hat{y}}(t; \theta)$  need to be learned
- ...Classical ODE integration methods are no longer viable

PINNs circumvent this issue by **using NN training for ODE integration**



# Training in PINNs

**In particular, we can apply a Lagrangian relaxation to the problem**

We relax the constraints in the previous formulation so that we obtain:

$$\begin{aligned}\mathcal{L}(y, \hat{y}, t, \theta) = & L(\hat{y}(t; \theta), y) \\ & + \lambda_{de}^T \|\dot{\hat{y}}(t; \theta) - f(\hat{y}(t; \theta), t)\|_2^2 \\ & + \lambda_{bc}^T \|\hat{y}(t_0; \theta) - y_0\|_2^2\end{aligned}$$

In optimization, this is called a **Lagrangian**

- Besides the original loss  $L$
- ...There is a penalty term linked to the ODE, with weights (multipliers)  $\lambda_{de}^T$
- ...And a penalty term linked to the initial value, with multipliers  $\lambda_{bc}^T$



# Training in PINNs

## The approach can be generalized

- In particular we can take into account both ODEs and PDEs
- ...And we can use different types of penalizers

## We just need to abstract a bit the formulation:

$$\begin{aligned}\mathcal{L}(y, \hat{y}, t, \theta) = & L(\hat{y}(t; \theta), y) \\ & + \lambda_{de}^T L_{de}(F(\hat{y}, t; \theta)) \\ & + \lambda_{bc}^T L_{bc}(B(\hat{y}, t; \theta))\end{aligned}$$

- Where  $F(y, t; \theta) = 0$  defines the original ordinary or partial DE
- ...And  $B(y, t; \theta) = 0$  defines the original initial or boundary conditions
- The  $L_{de}$  and  $L_{bc}$  terms can be L2 norms, but also other types of penalizer



# Training in PINNs

Then we train by:

- Sampling points  $\{t_i\}_{i=1}^n$  in the input space
- Choosing  $\theta$  so as to minimize the sum of Lagrangians

$$\operatorname{argmin}_{\theta} \sum_{i=1}^n \mathcal{L}(y, \hat{y}, t, \theta)$$

We can employ gradient descent, as usual

**Again, there is no need to use ODE/PDE integration at training time**

...Because *training is the integration process*

- In fact, it is possible to drop the data-based loss  $\mathbf{L}$  and the approach still works
- In such a case, PINNs can act as *approximate* ODE/PDE integrators





# No Free Lunch

**In the above description, it's easy to miss an important point**

Let's consider again the DE-based components in the Lagrangian:

$$L_{de}(F(\hat{y}, t, \theta)) \quad \text{which could be e.g.} \quad \|\dot{\hat{y}}(t; \theta) - f(\hat{y}(t; \theta), t)\|_2^2$$

- The penalizer contains **derivatives** (possibly partial)
- ...And it should provide a contribution for gradient descent

**This means that we need a way to compute the components of  $\dot{y}$**

...So that we obtain an expression that is **again differentiable in  $\theta$**

- This can be a bit tricky in practice!
- Viable approaches include symbolic differentiation (manual or automatic)
- ...Or partially numeric methods such as finite differences, etc.



# No Free Lunch

Moreover, assigning a value to the multipliers is not trivial

These the "weight" vectors  $\lambda_{de}$  and  $\lambda_{bc}$

$$\mathcal{L}(y, \hat{y}, t, \theta) = L(\hat{y}(t; \theta), y) + \lambda_{de}^T L_{de}(F(\hat{y}, t; \theta)) + \lambda_{bc}^T L_{bc}(B(\hat{y}, t; \theta))$$

Finding a good balance might be very tricky

- A good alternative might be using dual ascent

Finally, boundary conditions are **incorporated at training time**

...So, if they change, we need to **repeat training**

- In some contexts, this can be a major problem



# No Free Lunch

**Finally, one should be careful with the problem semantic**

Let's consider for a given input vector  $\mathbf{t}$  the constraint:

$$\|\dot{\hat{\mathbf{y}}}(\mathbf{t}; \boldsymbol{\theta}) - f(\hat{\mathbf{y}}(\mathbf{t}; \boldsymbol{\theta}), \mathbf{t})\|_2^2$$

**The constraint is enforced in a **soft** fashion**

...Meaning that it might be violated

- Proper weight calibration can help, but violations will still typically occur

**Even if we manage exact satisfaction**

...The constraint will hold only locally, for the specified  $\mathbf{t}$  values

- When we move away from the  $\mathbf{t}$  values considered at training time
- ...The NN may behave inconsistently with the underlying physics



## Some Remarks

**Let's conclude with some differences between mainstream PINNs and UDEs**

Unlike UDEs, PINNs need to **learn the involved physics**

- It might be necessary to use larger networks
- ...Because they will need to learn a more complex relation

The DE constraints are only **approximately satisfied**

- UDEs provide instead full guarantees
- ...But approximate satisfaction might be good if the DE is not fully reliable

PINNs do not rely on DE integration: they **are** an integration method

- This makes them faster than UDEs at inference and possibly training time
- ...But don't forget that changing the boundary conditions requires retraining!



## Some Remarks

### If you are looking for additional information

- There's a very well done PyTorch library for PINNs
- A well-known library is also available for JAX
- The PINN idea can be generalized, leading to Neural Operators
- ...Which map boundary conditions into integrated differential equations

